

Comparisons of different AI techniques on the game 2048

Written by Denis Legue (000513162), Robin Marchand, Gabriel, Monwar Mahtab Hawlader (000514679)

ULB
INFO-H410

Abstract

GitHub Repository —

<https://github.com/Somsro/infoH410>

Introduction

In the course INFO-H410 Techniques of Artificial Intelligence at ULB, students are required to apply and compare different AI techniques to the same problem. The game chosen is 2048. 2048 is a single-player game on a 4×4 grid where the player has to combine tiles by shifting them in one of the four possible directions: up, down, left and right. Each turn, a new tile with a value of 2 or 4 appears in an empty spot on the board. The player wins by creating a tile with the value of 2048, but can continue playing to reach higher values. The game ends when there are no more valid moves available.

The randomness of the games makes it a good candidate for testing the performance of different AI techniques. Three agents will be compared; Deep qlearner, Expectimax and an n-tuple network TD learning agent. The goal will be to evaluate their performance (the final score reached) but also the computational cost to better understand the trade-offs between planning and learned policies.

Problem Formulation

The game 2048 is a sequential decision making problem with uncertainty. The state of the game is represented by the configuration of the board. The configuration of the board is a 4×4 grid where each cell can be empty or contain a tile with a value that is a power of 2. At each time step, the agent observes the current board configuration and selects in which direction to shift the tiles in one of the four possible directions: up, down, left and right as long as there is at least one tile moving. Two tiles with the same value that collide during a shift will merge into a tile with the value of their sum. A new tile with a value of 2 or 4 will then appear in a random empty spot on the board which makes future states only partially controllable by the agent. The objective

of the agents is to achieve the highest score. The game ends when there are no more valid moves available. The fact that a new tile appears in a random empty spot on the board after each move makes the game stochastic and thus making it a good candidate for testing the performance of different AI techniques.

Formal Representation:

State Space (S): Values and grid positions of all tiles.

Actions (A): $\{Up, Down, Left, Right\}$, restricted to moves that modify the board.

Reward (r): Can be constructed based on several features (empty cells, monotonicity, and max tile placement).

Constraints: Terminal state is reached when no legal actions remain. Tabular Q-learning and exhaustive search are infeasible due to the massive state space.

Probabilities: Stochastic transitions $P(s' | s, a)$ dictate the game dynamics, preventing exact deterministic planning.

Methodologies

Three approaches will be implemented and compared; Expectimax, Deep Q-Learning and an n-tuple network TD learning agent. Each approach has its own field of application and is based on different principles. Expectimax is a search-based approach, deep Q-Learning is a deep reinforcement learning method that learns action values from experience meanwhile TD learning with function approximation is a lighter value-based alternative. Comparing them on the same problem will allow us to better understand the trade-offs between planning, learning, the computational cost as well as performance of each method. For reinforcement learning, a classic approach is to use a Q-learning algorithm with a tabular representation of the state-action values. However, due to the large state space of the game, this approach is not feasible, in fact, the number of possible board configurations is estimated to be around 12^{16} , if we consider that the maximum tile value is 2048 (which is 2^{11}). The following sections will describe how this issue is addressed in the case of the Deep Q-Learning and the TD learning agent. For each of those approaches, the state is encoded as a list of 16 integers, each representing the exponent k such that the tile value in that cell is 2^k . (For example, 16 is encoded as 4 since $2^4 = 16$ with the particularity that 0 represents an empty cell). The action space is the same for all agents and

consists of the 4 possible moves: up, down, left and right encoded as integer values from 0 to 3.

Approach: Expectimax

Expectimax is a search algorithm that extends the minimax algorithm to handle randomness in the environment. **Logic:** The agent constructs a search tree where nodes alternate between decision nodes (agent's turn) and chance nodes (random tile placement). It evaluates the expected utility of actions by averaging over possible outcomes at chance nodes, using a given evaluation function to estimate the leaf nodes. **Evaluation Function:** Many heuristics exist to evaluate the board state, the implemented one is a combination of several features:

- Raw score, which is the sum of the values of every tile on the board
- Number of empty cells
- Monotonicity, which measures how values are ordered on the board
- Smoothness, which measures how similar neighboring tiles are
- Merge potential, which estimates the possibility of merging tiles in the next moves
- Corner bonus/penalty, which rewards/penalizes the presence/absence of the largest tile in a corner

Depth and Computational Cost: The depth of the search tree is a critical parameter that affects both performance and computational cost. A deeper tree allows for better foresight and potentially higher scores, but it also increases the number of nodes exponentially, leading to longer computation times. Tuning this parameter is essential to find a balance between performance and efficiency, especially given the stochastic nature of the game. Implementing pruning can lead to significant reductions in the number of nodes evaluated without sacrificing performances, thus considerably reducing the computational cost. The pruning method implemented is a simple threshold-based pruning that discards branches of the search tree that have low chance of happening.

Approach: Deep Q-Learning (DQN)

Logic: Approximate the action value function (Q-value) using a deep neural network to bypass the exhaustive exploration of the state space. The agent learns directly from experience via stochastic gradient descent on (s, a, r, s') tuples, guided by a shaped reward system reflecting 2048 heuristics (tile merges, empty cells, monotonicity, and corner penalties).

Assumption: The environment features a massive state space and stochastic transitions, rendering tabular methods infeasible. While this model-free approach is good at capturing complex patterns, it assumes access to high computational power and massive amounts of training data (very long training sessions) to converge.

Approach: N-tuple network TD(0) learning agent

Temporal-Difference learning with value function approximation provides a lighter alternative to Deep Q-Learning.

Logic: Instead of learning action values directly, this approach estimates the value of board after-states and updates this estimate by bootstrapping from successor after-states. In this implementation, the after-states represent the move without placing a new tile yet in order to estimate the value function. An error signal is computed based on the difference between the current value estimate and the observed reward plus the value of the next after-state. The value function is approximated using an N-tuple network composed of lookup tables defined over several tile patterns and their symmetric variants.

Algorithm 1: Training procedure of the TD agent

Require: Number of episodes N

```

1: Initialize TD agent with N-tuple network  $V\theta$ 
2: for episode = 1 to  $N$  do
3:   Reset the environment
4:   Initialize empty trajectory path
5:   done  $\leftarrow$  false
6:   while not done do
7:     Select an action using  $\epsilon$ -greedy exploration
8:     Clone the environment and apply the action to
       obtain the after-state and reward
9:     Execute the action in the real environment
10:    Store the after-state, reward, and terminal flag in
       the path
11:    end while
12:     $target \leftarrow 0$ 
13:    for each stored after-state  $s$  in reverse order do
14:       $v \leftarrow V\theta(s)$ 
15:       $\delta \leftarrow target - v$ 
16:      Update the weights of  $V\theta$  using  $\delta$ 
17:      Update  $target \leftarrow r + V\theta(s)$ 
18:    end for
19:    Update  $\epsilon \leftarrow \max(\epsilon_{\min}, \epsilon \cdot \lambda)$ 
20:  end for
21: return trained agent

```

References